Why-Agent

The agent that can explain why

OAH 2026 · Track 02 · Team MENfigiTAL

Problem

Agents act, but can't show **why**. Founders can't trust an action they can't audit — especially money moves.

What it does

Every action returns a **decision card**: rules applied · evidence with sources · alternatives scored · 3 reviewers (2/3 quorum, silence = NO) · confidence · what would refute it · past precedents. Human says Yes / No / Later.

Architecture

Deterministic core (rules P1–P12, P7 hard-blocks the agent from signing/sending funds)
Evidence: Zetaris REST → local JSON fallback
Memory: Meterless H-MEM engine (remember · recall with provenance · trust-ledger audit)
Narration: optional LLM, never changes the decision

Meterless engine in use

Each decision is stored in H-MEM. New questions recall precedents via hybrid retrieval; the card shows the memory id and its ledger history (create→read). Zero keys, runs locally.

Try it

github.com/mdm-delo/why-agent

`python3 server.py` → localhost:8765

`python3 agent.py --demo`